

NAVEEN KUMAR S

Coimbatore, Tamil Nadu | +91 70927 03991 | nav.cs@outlook.com
linkedin.com/in/nav-cs | github.com/naveendha1710

CAREER OBJECTIVE

AI & Tech enthusiast building hands-on across agentic systems, LLM infra, and applied AI — from a production ERP (50,000+ assets tracked) to LLM gateways and model quantization pipelines. Looking to grow into AI Engineer / AI R&D roles where I can keep going deeper into how these systems actually work.

EDUCATION

B. Sc Artificial Intelligence & Machine Learning | CGPA – 7.46 2024 – 2027
Rathinam global deemed to be university, Coimbatore
HSC – 83.6 | SSLC - 72.4

INTERNSHIP EXPERIENCE

- AI Developer** — *Rathinam Group of Institutions* Sep 2025 – Present
- Developed an institution wide ERP System as an AI developer, currently live with 100+ User and active in production use
 - Built modules for ticketing and asset management, enabling tracking of 50,000+ labeled assets across the institution
- Software Tester** — *SMARTMATE Systems* Jan 2025 – Aug 2025
- Identified a reported 80+ UI/UX bugs across the application, improving user experience and software stability
 - Conducted cross-device testing to ensure consistent functionality and interface quality

PROJECTS

- ERP System** | Full Stack + AI | *React.js, Node.js, TypeScript, PostgreSQL, Supabase, MCP*
- Built a production-grade ERP handling Tenant, Asset management, Helpdesk and 50,000+ tracked assets
 - Implemented agentic AI workflows for automated task routing and management across institutional departments
- Semantic Router & LLM Caching Gateway** | *Python, FastAPI, Redis, Docker*
- Built an LLM gateway with semantic caching (vector similarity, Redis-backed) and complexity-based routing across cloud and local providers, reducing repeated API calls.
 - Added quality-gated caching, and provider failover; benchmarking, and load testing, backed by 145 automated tests.
- LLM Quantization Pipeline** | *Python, Hugging Face, optimum-quanto*
- Built a pure-Python pipeline to download, quantize (INT8/INT4), benchmark, and run LLMs (Qwen2.5, Llama 3) locally on CPU, without relying on pre-quantized downloads
 - Designed modular CLI scripts to benchmark size, load time, and tokens/s across BF16, INT8, and INT4 quantization levels, plus an interactive chat mode

TECHNICAL SKILLS

Programming Languages: Python, C++, Go, SQL | **With AI:** React.js, TypeScript, Node.js, Tailwind CSS, REST APIs, SMTP
AI / ML & Automation: OpenRouter, N8N, OPAL, Ollama, vLLM, MCP, CopilotAgents, OpenClaw, NvidiaBuild, RAG
Databases: PostgreSQL, Qdrant, Redis
Tools & Platforms: Git, Ubuntu (Desktop/Server), Docker, Supabase, Postman MCP, ClaudeCode CLI, MobaXterm, Orange

HACKATHONS & COMPETITIONS

- BNB Chain x YZi Labs Hack Series, Bengaluru (Feb 2026) — Shortlisted to Finals (Top 50) | Web3 Track
- Smart India Hackathon (SIH) — Proposed DWLR-based smart water conservation system for wells

CERTIFICATIONS

- Docker Foundations Professional Certificate — Docker & LinkedIn
- Building with the Claude API, Claude 101 & Model Context Protocol — Anthropic
- Generative AI Engineering — IBM (Coursera)

LANGUAGES

English, Tamil